

When Do Structural Primers Help LLMs Reason over Graphs?

Baseline Accuracy and Answer Leakage

Gal Barak

github.com/GalBrk

Nitzan Zacharia

github.com/NitzanZacharia

Inbal Moryles

github.com/iinbal

Abstract

Prepending a primer of structural facts, such as each node’s degree, is a natural aid for a language model answering questions about a graph given as text. We ask when such primers help, and why. We evaluate Qwen3-1.7B and Qwen3-4B, with and without thinking, on six GraphQA tasks over 40-node Erdős–Rényi graphs (84,000 generations). A solver that never sees the edges separates primers that state the answer from primers that supply side information. On node-degree queries, a primer that states the answer adds about 16 points where accuracy without it lies between 0.25 and 0.75, and it costs Qwen3-4B accuracy where that model is already correct. The responses indicate why: the primer changes the model’s procedure. Qwen3-4B mostly retrieves the stated degree instead of counting; it gains 27 points where its count fails and loses 12 where its count is 99% accurate. With thinking, Qwen3-1.7B instead checks its count against the stated degree and gains 20 points. On edge queries, primers act on false “yes” answers: for Qwen3-1.7B, the combined primer lowers the false-alarm rate by 19 points, and a preamble that only restates the node set raises it by 18. Primer length has no consistent effect.

1 Introduction

Fatemi et al. (2024) show that the accuracy of a large language model (LLM) on graph questions depends strongly on how the graph is written as text. A natural next step is to add derived structure to that text. We study *primers*: short preambles, placed before the graph encoding, that state each node’s degree, clustering coefficient or random-walk return probability, or the graph’s number of connected components. Graph neural networks take these statistics as input features (Dwivedi et al., 2022; Rampášek et al., 2022); here a pretrained LLM reads them as sentences.

Our research question is how primers change an LLM’s answers on graph tasks, and when they

help rather than harm. We divide it into three parts.

RQ1: how much of a primer’s effect comes from stating the answer, and how much from side information that does not determine it? **RQ2:** what decides whether a primer helps or harms a given model on a given task? **RQ3:** what does a primer change in the procedure the model follows?

Two controls make these questions answerable. A graph-blind solver reads the primer but never the edge list; where it is correct, the primer states the answer (*answer leakage*), and a gain may reflect retrieval rather than reasoning (Geirhos et al., 2020). Because added text can itself degrade reasoning (Shi et al., 2023; Levy et al., 2024), a preamble that only restates the node set and primers of 37 to 4,691 characters test whether added text alone changes accuracy. We apply both controls to 84,000 generations and analyse the responses as well as their scores.

The results suggest one account (Section 5): a primer that states the answer changes the procedure the model follows, and its effect depends on how accurate the new procedure is relative to the count it replaces. The account covers three observations: on node_degree such effects follow the accuracy without a primer (Section 4.1), one primer both helps and harms the same model on the same task (Section 4.2), and a stated answer helps only a model that can execute the procedure it implies (Section 4.3). On edge queries, side information instead changes how often the model answers “yes” (Section 4.4).

2 Related Work

GraphQA (Fatemi et al., 2024), NLGraph (Wang et al., 2023), GPT4Graph (Guo et al., 2023) and GraphArena (Tang et al., 2025) evaluate LLMs on graph problems and vary the encoding, the prompt and the task. Sanford et al. (2024) show that four of our six tasks can be solved by a single transformer layer, yet prompted LLMs trail trained transformers

Condition	States	Chars	Carries
none	no primer	–	–
components	component count	37	–
degree	degree	891	deg., count
clustering	clustering coefficient	1,629	–
rwse	return prob., 2 and 3 steps	2,949	–
all	the three rows above	4,691	deg., count
filler	“Node i is simply present”	1,829	–

Table 1: The seven conditions. *States*: what each primer gives per node (components: for the graph). *Chars*: mean length. *Carries*: tasks with variable gold on which the graph-blind solver scores at least 0.85 (Appendix Table 5): node_degree (deg.) and edge_count (count).

on them. Structure has been added to the input as an instruction to build the graph first (Wang et al., 2023), as motif statistics (Das et al., 2024) and as learned soft tokens (Perozzi et al., 2024). Huang et al. (2024) find that such text helps node classification through label-relevant phrases rather than through structure. We make the analogous separation for every task and condition, between text that determines the answer and text that does not. In the taxonomy of Rampásek et al. (2022), our primers state local structural encodings (degree, triangle-based clustering and random-walk return probability; Dwivedi et al., 2022) and one global encoding, the number of components.

The graph-blind solver is a partial-input baseline, as hypothesis-only models are for natural language inference (Gururangan et al., 2018; Poliak et al., 2018), applied here to an intervention rather than to a dataset. Accuracy also depends on where information sits in the context (Liu et al., 2024) and on the order of a graph description (Ge et al., 2025).

3 Method

3.1 Primers and the graph-blind solver

A prompt consists of a primer, the graph in the incident encoding of Fatemi et al. (2024), and the question (Appendix A). Seven conditions share one renderer (Table 1). The five *content* primers state graph statistics, one sentence per node in id order, except components, which is a single sentence. filler writes “Node i is simply present within the graph G ” for each node: it restates the node set, which the encoding already lists, at a length 12% above that of clustering.

The graph-blind solver reads the rendered primer and the graph’s node and edge counts, never the edge list. It applies exact rules where they exist

and heuristic or fitted rules otherwise (fitted on separate graphs), so its accuracy is a lower bound on what the primer and the two counts determine. We call a primer *answer-carrying* on a task when the solver scores at least 0.85 there, and *side information* otherwise. Any threshold between 0.75 and 1.00 gives this split: on the four tasks whose gold answer varies, degree and all score 1.00 on node_degree and edge_count, and every other primer scores at most 0.75 (Appendix Table 5).

3.2 Corpus, models and measurement

Graphs are Erdős–Rényi $G(n, p)$ graphs (Erdős and Rényi, 1959; Gilbert, 1959) with $n=40$ nodes at $p \in \{.10, .20, .35, .50\}$, 100 per density. The same graphs serve all six GraphQA tasks, and node_degree, connected_nodes and edge_existence query the same node. node_degree and edge_existence are extended to $p \in \{.65, .75, .85\}$. Four arms answer every prompt: Qwen3-1.7B and Qwen3-4B (Qwen Team, 2025b) with thinking mode off (qwen3-1.7b, qwen3-4b) and on (-think), 84,000 generations. Decoding is greedy, with a budget of 8,192 new tokens (2,048 for the plain arms in the extension). For qwen3-1.7b on node_degree we use three further runs: 400 graphs per density (the main sweep’s 100 and 300 new), 1,600 graphs drawn with fresh seeds, and a grid that holds the mean degree at 8 or 16 while n ranges over $\{20, 40, 80, 160\}$.

Every task is scored by exact match; on connected_nodes, set-F1 moves about a sixth as much as exact match and reverses two of the five contrasts larger than 5 points. On edge_existence the share of queried pairs that are edges rises from 10% to 85% with density, so we also report the hit rate (“yes” to an edge), the false-alarm rate (“yes” to a non-edge) and balanced accuracy. Each effect is a paired difference against none on the same graphs, in percentage points. A pair is dropped when either generation reaches the token budget (rates in Appendix Table 6). A *cell* is one (arm, task, density, content primer) combination; we keep the 393 cells that retain at least 50 pairs, and the 47 dropped are edge_count cells. node_count and cycle_check have a constant gold answer at $n=40$, so we exclude them from all summaries (Appendix B). Tests are exact McNemar (McNemar, 1947) with Benjamini–Hochberg correction (Benjamini and Hochberg, 1995) within each (arm, task), and intervals are 95% bootstrap




Figure 1: Effect of each content primer against the cell’s accuracy without it. Triangles: answer-carrying; circles: side information; filled: $q < .05$ within each (arm, task). The line joins the seven densities of degree for qwen3-4b on node_degree; shading marks baselines 0.25–0.75.

intervals over graphs (Efron, 1979).

4 Results

4.1 Primer effects follow baseline accuracy

Figure 1 plots each cell’s effect against its accuracy without a primer, its *baseline*. Answer-carrying primers gain most at intermediate baselines: +16.0 and +16.7 points at baselines of 0.25–0.50 and 0.50–0.75 (all 18 such cells are node_degree cells), against +7.8 below 0.25 and –5.7 between 0.75 and 0.90. Above 0.90 they cost 1.4 points on average, and every such cell of qwen3-4b loses accuracy (–1 to –15). Side information follows the same shape at a smaller scale: +2.3 and +3.9 in the two middle bands, –1.8 below 0.25 and –0.8 above 0.90. The pattern is not regression to the mean: binning each cell on half its graphs and measuring the effect on the other half reproduces every band to within about 3 points.

Thinking raises the median baseline from 59.0 to 94.9 on the 1.7B model and from 87.0 to 100.0 on the 4B model. It empties the range below 0.25, which holds 28 and 20 cells without thinking, and leaves the 1.7B model a similar share between 0.25 and 0.75 (40 of 103 cells, then 35 of 90) but the 4B model none (10 of 110 before). Pooled over $p \leq .50$, no primer moves qwen3-4b-think by more than 2.1 points on node_degree, edge_existence or connected_nodes (Table 2).

Table 2 breaks the effects down by task. Answer-carrying primers raise node_degree on both 1.7B arms (+7.5 to +11.7) and lower it on qwen3-4b (–6.5, –8.0), and degree raises qwen3-4b

on edge_count by 27.7. As side information, all helps qwen3-1.7b on edge_existence (+11.5), clustering is the only primer that helps qwen3-4b there (+4.0), and filler produces the largest costs (–15.7, –12.5).

4.2 The sign follows the procedure the primer replaces

For qwen3-4b on node_degree, the degree primer states the queried degree and changes accuracy by –6, –1, –7 and –12 points at $p=.10$ to .50, by –6 at .65, by +7 at .75 and by +27 at .85 (Figure 1, line). The losses are not a cost of added text: at $p=.50$ the primer turns 13 correct answers into errors and corrects one, no generation truncates, and filler, which is longer, costs 2 points.

The responses show a change of procedure. We classify a node_degree response as *retrieval* when it states its answer before restating the queried node’s neighbour list, or without restating it, and as *enumeration* when it lists the neighbours one per line (Appendix Table 4). Without a primer there is no stated degree to retrieve, so the baseline of qwen3-4b is the accuracy of its own count: 99–100% up to $p=.50$, then 82%, 50% and 33%. Under degree it retrieves in 47–97% of responses, at 57–97% accuracy. Retrieval is less accurate than the count up to $p=.65$ and more accurate at .75 and .85, which matches the sign of the effect at all seven densities. Enumeration falls from 85–100% of responses to at most 8% at $p=.35$ to .75. Under all, which states the degree among two other statistics, the model retrieves in 27–86% of responses at 40–94% accuracy, its other responses are 43%, 4% and 2% accurate from $p=.65$, and the primer lowers accuracy at every density, by up to 37 points.

The thinking arm uses the same primer differently: under degree, qwen3-1.7b-think retrieves in at most 2% of responses and, from $p=.35$ on, enumerates the neighbours in 76–98%. It reports a discrepancy between its count and the primer in 26% of responses up to $p=.50$ and in 49% beyond, whereas without a primer 1% and 4% of responses report any discrepancy, and it answers correctly in about two thirds of the responses that report one. The primer adds 19.9 points [16.3, 23.6] over the seven densities and 31.6 points above $p=.50$, without the loss at high baselines (+5.2 and –1.0 at $p=.10$ and .20). qwen3-1.7b, finally, rarely retrieves (at most 4% of responses), yet gains 10

	qwen3-1.7b				qwen3-1.7b-think				qwen3-4b				qwen3-4b-think			
	exist.	deg.	conn.	count	exist.	deg.	conn.	count	exist.	deg.	conn.	count	exist.	deg.	conn.	count
none (acc.)	70.5	60.5	72.5	1.8	99.2	76.8	79.7	86.3	90.5	99.2	86.0	2.0	100.0	99.7	98.7	97.6
pairs	399	399	399	211	397	376	390	13	398	400	400	393	393	380	373	34
filler	-15.7	-0.5	-6.2	-0.4	-0.8	-3.5	+2.0	+3.4	-7.0	-0.5	-12.5	+1.2	+0.0	+0.0	+0.8	-2.0
components	-8.2	+1.8	+0.8	-1.4	-0.8	-1.0	+2.8	+7.4	-2.8	+0.2	+9.3	-1.3	+0.0	-0.3	+0.5	+1.7
clustering	+3.8	+3.0	+1.3	-0.4	-1.0	+1.0	+1.5	-8.7	+4.0	-0.5	+1.0	-1.3	+0.0	-0.3	+0.0	+1.7
rwse	+5.3	+4.0	+1.0	-0.8	-1.3	+2.5	-0.3	-4.0	-10.0	-2.3	+0.0	-1.5	+0.0	-2.1	+0.0	+0.0
degree	+4.0	+7.5	+3.5	+4.0	-2.3	+11.7	+7.4	-61.5	-4.3	-6.5	-3.0	+27.7	+0.0	+0.0	+0.0	-23.5
all	+11.5	+10.3	+2.0	+2.4	-2.0	+11.6	+6.1	-8.3	-9.8	-8.0	-2.7	+3.8	+0.0	-0.8	-0.5	-41.5

Table 2: Pooled effects against none (points) over $p \leq .50$. Tasks: edge_existence (exist.), node_degree (deg.), connected_nodes (conn.), edge_count (count). First row: accuracy under none over terminated generations; *pairs*: fewest untruncated pairs. **Bold**: $q < .05$ (exact McNemar, Benjamini–Hochberg over each column’s six contrasts); columns with *italic* pair counts (below 100) are never bolded, since their pairs are selected on termination.




Figure 2: qwen3-1.7b on edge_existence: (a) balanced accuracy and (b) median new tokens by density. Without a primer the model answers “yes” to almost every pair from $p=.65$; under filler, from $p=.35$.

and 18 points at $p=.20$ and $.35$, where its count is partly right (80% and 41%), and stays within 2 points from $p=.50$, where its baseline is 5–30%.

4.3 A stated answer helps only through a procedure the model can execute

On edge_count, degree states every degree, so the answer is half their sum. Without a primer qwen3-4b already sets out to sum the degrees (96–100% of responses) but is exact on only 2% of graphs. Given the degrees it is exact on 30% of graphs: +23, +30, +24 and +34 points across the four densities, with 114 items corrected and 5 broken. qwen3-1.7b takes up the same procedure (from 7–67% of responses to 99–100%) but rarely completes the forty-term sum correctly: 4.7% of its answers are exact, a gain of 4.0 points that does not survive correction, although its mean absolute error falls from 136 to 40 edges.

4.4 Side information shifts false alarms and helps where counting fails

On edge_existence the plain arms answer “yes” to 97–100% of true edges under every condition,

and 94–99% of their errors are false alarms, so primers change accuracy almost entirely through the false-alarm rate. Pooled over the seven densities, all lowers the false-alarm rate of qwen3-1.7b from 0.51 to 0.33 (−18.6 points [−24.0, −13.1]), and degree, rwse and clustering lower it by 6 to 9 points; filler raises it to 0.69 (+18.3 [13.9, 22.7]) and components to 0.61. For qwen3-4b, clustering is the only primer that lowers the false-alarm rate (0.16 to 0.08) and the only one that raises accuracy on the task (+4.6 [2.3, 6.9]); the other five conditions raise false alarms by 7 to 14 points.

At high density qwen3-1.7b defaults to “yes” (Figure 2). Without a primer it answers “yes” to 98–99% of pairs from $p=.65$, in median responses of 48–120 tokens, at a balanced accuracy of 0.52–0.53. Primers that state per-node values keep the responses at 259–404 tokens, and the two that state degrees restore part of the balanced accuracy (degree +9.5 [4.5, 15.0], all +14.9 [8.7, 21.4]). filler brings the default forward to $p=.35$, where the model answers “yes” to 97% of pairs against 65% without a primer.

Side information also helps where counting fails. On node_degree above $p=.50$, where qwen3-4b counts correctly on 33–82% of items, clustering adds 11.3 points [6.7, 16.3] (46 items corrected, 12 broken) and rwse adds 7.7, whereas filler adds 2.0 (not significant). These primers state no degree, so the answer still comes from the model’s own count, and at these densities the printed coefficients barely vary within a graph (standard deviation at most 0.02), so the gain does not rest on node-specific values. The solver scores clustering as it scores none, yet the primer also helps qwen3-1.7b: its gain of +3.0 in the main sweep replicates in

the three dedicated runs, at +4.0, +4.2 and +6.2 points on the 300 new graphs per density, the fresh seeds and the fixed-degree grid ($p \leq .005$ in each).

Some gains need no graph-specific content: components raises qwen3-4b on connected_nodes by 9 to 11 points at $p \geq .20$, where every graph is connected and the primer is the same 37-character sentence for all of them, and by +7 at $p=.10$, where it varies.

4.5 No consistent gain from bundling or length

all states every value that degree, clustering and rwse state, yet its effect exceeds the mean of theirs by only 0.2 points on average ($[-1.1, 1.6]$ over 54 (arm, task, density) combinations). It can also lose what each part gains: on qwen3-4b/node_degree above $p=.50$, degree, clustering and rwse add 9.3, 11.3 and 7.7 points, whereas all costs 21.3 (Section 4.2).

Primer length does not order the effects consistently. Regressing each (arm, task)’s effects on the lengths of the five content primers gives slopes from -6.5 to $+3.4$ points per 1,000 characters, positive in 7 of 8 fits for the 1.7B model and in none of 8 for the 4B model; without all, the longest primer, 3 and 1 of 8 are positive. filler lies 18.3 and 13.8 points below the fitted line where it harms most. On node_count, which the summaries exclude, qwen3-1.7b answers 39 instead of 40 on 68.5% of items without a primer, 1.2% under clustering and 66.5% under the longer rwse (Appendix Table 3).

5 Discussion and Conclusion

The results suggest one account. A primer that states the answer changes the procedure the model follows, and its effect depends on how accurate the new procedure is relative to the count it replaces. Without a primer the model must count, so its baseline measures the count a primer can displace, which is why these effects follow it. For qwen3-4b the new procedure is retrieval, which beats its count only from $p=.75$; all makes retrieval less accurate and its other responses fail at high density, so it loses even where its parts help. The 1.7B arms keep counting, and the thinking arm uses the primer as a check. A new procedure must also be executable: given every degree, qwen3-4b completes the sum exactly on 30% of graphs and qwen3-1.7b on 5%, a failure of multi-step compo-

sition of the kind Dziri et al. (2023) describe. Side information that does not state the answer acts on another channel: on edge queries it changes how often the model answers “yes”.

Two results bear on prior work. A single transformer layer can solve node_degree and edge_count (Sanford et al., 2024), the two tasks on which a primer states the answer, yet given the answer qwen3-4b is right on 87% of node_degree items at $p=.50$ and on 30% of edge_count graphs: for a prompted model the difficulty appears to lie in the procedure, not in the information. Added text acts through what it states: restating the node set raises false “yes” answers, consistent with the distraction Shi et al. (2023) report, while the length slopes change sign between models, unlike the consistent cost Levy et al. (2024) report.

In answer to our questions: on node_degree a stated answer is worth about 16 points at intermediate baselines, and side information 2 to 4 points across tasks (RQ1); the sign of a stated answer’s effect follows the accuracy of the count the primer displaces (RQ2); and a primer changes the model’s procedure, to retrieval or verification on node_degree and, on edge queries, to a different rate of “yes” answers (RQ3). Evaluations of graph prompts should report each cell’s baseline accuracy, a graph-blind solver’s score and, for yes/no tasks, the false-alarm rate.

Limitations

We evaluate two model sizes from one family, with one encoding, template and primer position. Decoding is greedy, whereas the Qwen3 model card (Qwen Team, 2025a) recommends sampling in thinking mode, so the thinking arms’ truncation on edge_count (81% and 62%) may be specific to this protocol. A cell of 100 graphs detects paired effects of about 8 points, so the side-information results rest on averages over cells. Procedures are classified by surface patterns, and a model may follow a procedure it does not state. The length comparison varies length only across primers whose content also differs.

Ethics Statement

This work evaluates open-weight models on synthetic graphs, with no human subjects, personal data or annotation. Its main cost is GPU compute: about 217,000 generations over the project, 98,400 of them analysed here.

References

- Yoav Benjamini and Yosef Hochberg. 1995. [Controlling the false discovery rate: A practical and powerful approach to multiple testing](#). *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1):289–300.
- Debarati Das, Ishaan Gupta, Jaideep Srivastava, and Dongyeop Kang. 2024. [Which modality should I use – text, motif, or image? : Understanding graphs with large language models](#). In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 503–519.
- Vijay Prakash Dwivedi, Anh Tuan Luu, Thomas Laurent, Yoshua Bengio, and Xavier Bresson. 2022. [Graph neural networks with learnable structural and positional representations](#). In *International Conference on Learning Representations*.
- Nouha Dziri, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jiang, Bill Yuchen Lin, Sean Welleck, Peter West, Chandra Bhagavatula, Ronan Le Bras, Jena D. Hwang, Soumya Sanyal, Xiang Ren, Allyson Ettinger, Zaid Harchaoui, and Yejin Choi. 2023. [Faith and fate: Limits of transformers on compositionality](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 70293–70332.
- Bradley Efron. 1979. [Bootstrap methods: Another look at the jackknife](#). *The Annals of Statistics*, 7(1):1–26.
- Paul Erdős and Alfréd Rényi. 1959. [On random graphs I](#). *Publicationes Mathematicae Debrecen*, 6(3–4):290–297.
- Bahare Fatemi, Jonathan Halcrow, and Bryan Perozzi. 2024. [Talk like a graph: Encoding graphs for large language models](#). In *International Conference on Learning Representations*.
- Yuyao Ge, Shenghua Liu, Baolong Bi, Yiwei Wang, Lingrui Mei, Wenjie Feng, Lizhe Chen, and Xueqi Cheng. 2025. [Can graph descriptive order affect solving graph problems with LLMs?](#) In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6404–6420.
- Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. 2020. [Shortcut learning in deep neural networks](#). *Nature Machine Intelligence*, 2(11):665–673.
- E. N. Gilbert. 1959. [Random graphs](#). *The Annals of Mathematical Statistics*, 30(4):1141–1144.
- Jiayan Guo, Lun Du, Hengyu Liu, Mengyu Zhou, Xinyi He, and Shi Han. 2023. [GPT4Graph: Can large language models understand graph structured data? An empirical evaluation and benchmarking](#). *arXiv preprint arXiv:2305.15066*.
- Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. 2018. [Annotation artifacts in natural language inference data](#). In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 107–112.
- Jin Huang, Xingjian Zhang, Qiaozhu Mei, and Jiaqi Ma. 2024. [Can LLMs effectively leverage graph structural information through prompts, and why?](#) *Transactions on Machine Learning Research*.
- Mosh Levy, Alon Jacoby, and Yoav Goldberg. 2024. [Same task, more tokens: the impact of input length on the reasoning performance of large language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15339–15353.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. [Lost in the middle: How language models use long contexts](#). *Transactions of the Association for Computational Linguistics*, 12:157–173.
- Quinn McNemar. 1947. [Note on the sampling error of the difference between correlated proportions or percentages](#). *Psychometrika*, 12(2):153–157.
- Bryan Perozzi, Bahare Fatemi, Dustin Zelle, Anton Tsitsulin, Mehran Kazemi, Rami Al-Rfou, and Jonathan Halcrow. 2024. [Let your graph do the talking: Encoding structured data for LLMs](#). *arXiv preprint arXiv:2402.05862*.
- Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme. 2018. [Hypothesis only baselines in natural language inference](#). In *Proceedings of the Seventh Joint Conference on Lexical and Computational Semantics*, pages 180–191.
- Qwen Team. 2025a. Qwen3-4B model card. <https://huggingface.co/Qwen/Qwen3-4B>.
- Qwen Team. 2025b. [Qwen3 technical report](#). *arXiv preprint arXiv:2505.09388*.
- Ladislav Rampášek, Michael Galkin, Vijay Prakash Dwivedi, Anh Tuan Luu, Guy Wolf, and Dominique Beaini. 2022. [Recipe for a general, powerful, scalable graph transformer](#). In *Advances in Neural Information Processing Systems*, volume 35, pages 14501–14515.
- Clayton Sanford, Bahare Fatemi, Ethan Hall, Anton Tsitsulin, Mehran Kazemi, Jonathan Halcrow, Bryan Perozzi, and Vahab Mirrokni. 2024. [Understanding transformer reasoning capabilities via graph algorithms](#). In *Advances in Neural Information Processing Systems*, volume 37, pages 78320–78370.

Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Schärli, and Denny Zhou. 2023. [Large language models can be easily distracted by irrelevant context](#). In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 31210–31227.

Jianheng Tang, Qifan Zhang, Yuhua Li, Nuo Chen, and Jia Li. 2025. [GraphArena: Evaluating and exploring large language models on graph computation](#). In *International Conference on Learning Representations*.

Heng Wang, Shangbin Feng, Tianxing He, Zhaoxuan Tan, Xiaochuang Han, and Yulia Tsvetkov. 2023. [Can language models solve graph problems in natural language?](#) In *Advances in Neural Information Processing Systems*, volume 36, pages 30840–30861.

A Example prompt

A clustering-condition node_degree prompt at $p=0.10$ (elisions marked ...; line breaks added for width). Other conditions replace the first paragraph; none omits it.

Node 0 has clustering coefficient 0.20. Node 1 has clustering coefficient 0.00. ... Node 39 has clustering coefficient 0.33.

G describes a graph among nodes 0, 1, 2, ..., 38, and 39.

In this graph:

Node 0 is connected to nodes 2, 6, 8, 12, 35.

Node 1 is connected to nodes 5, 20.

...

Node 39 is connected to nodes 15, 16, 18, 20.

Q: What is the degree of node 16?

A:

B Constant-gold tasks

Arm	none	filler	comp.	clust.	rwse	degree	all
node_count							
1.7B	31.5	+39.8	+21.2	+67.3	+2.0	+6.8	+47.0
1.7B-think	98.6	+1.4	+0.9	+1.4	+1.4	+0.8	+1.4
4B	90.2	+9.3	+9.0	+8.8	+9.3	+8.0	+9.8
4B-think	99.5	+0.5	+0.0	+0.5	+0.5	+0.5	+0.5
cycle_check							
1.7B	99.7	+0.3	-0.6	+0.3	-7.5	-12.6	-1.1
1.7B-think	100.0	+0.0	+0.0	+0.0	+0.0	-0.7	+0.0
4B	99.7	-0.3	-1.3	+0.3	+0.3	-1.5	-0.5
4B-think	100.0	+0.0	+0.0	+0.0	+0.0	-0.8	+0.0

Table 3: The two tasks with constant gold at $n=40$, pooled over $p \leq .50$: accuracy under none (%) and paired effects against it (points). **Bold**: as in Table 2.

C Procedure on node_degree

p	.10	.20	.35	.50	.65	.75	.85
qwen3-4b							
acc., none	100	99	99	99	82	50	33
enumerates	1	34	95	100	94	85	44
acc., degree	94	98	92	87	76	57	60
retrieves	47	72	92	88	96	97	97
acc. of retrievals	89	97	92	85	77	57	62
enumerates	2	14	6	8	2	1	0
acc., all	93	96	92	84	45	31	25
retrieves	86	68	31	27	63	74	60
acc. of retrievals	92	94	77	63	46	41	40
acc. of the rest	100	100	99	92	43	4	2
qwen3-1.7b-think							
acc., none	93	95	60	60	52	47	42
acc., degree	98	94	84	82	78	76	87
enumerates	2	37	76	91	90	98	96
states a discrepancy	8	14	43	40	60	47	40
qwen3-1.7b							
acc., none	91	80	41	30	8	16	5
acc., degree	94	90	59	29	8	15	7
retrieves	4	2	0	0	0	0	4

Table 4: Procedure on node_degree by density (%), terminated generations). *Retrieves*: the response states its answer before restating the queried node’s neighbour list, or without restating it. *Enumerates*: it lists the neighbours one per line. *The rest*: responses that do not retrieve. *States a discrepancy*: it reports that its count and the primer disagree. Without a primer every arm counts, so the none rows give the accuracy of counting.

D Solver scores and truncation

Task	none	comp.	clust.	rwse	degree	filler	all
node_count	1.00	1.00	1.00	1.00	1.00	1.00	1.00
edge_count	.030	.030	.104	.030	1.00	.030	1.00
node_degree	.144	.144	.144	.210	1.00	.144	1.00
connected_nodes	.000	.000	.000	.001	.001	.000	.008
edge_existence	.735	.735	.727	.554	.746	.735	.741
cycle_check	1.00	1.00	1.00	1.00	1.00	1.00	1.00

Table 5: Accuracy of the graph-blind solver at $n=40$, averaged over seven densities (300 test graphs per density; fitted rules use separate graphs). It reads the rendered primer and the node and edge counts, never the edge list. **Bold**: answer-carrying (score ≥ 0.85) on a task whose gold varies. On node_count and cycle_check every score is 1.00 because the gold answer is constant at $n=40$.

Task	1.7B	1.7B-think	4B	4B-think
edge_count	22.4%	81.0%	0.6%	62.2%
cycle_check	6.1%	6.5%	1.0%	1.9%
node_degree	0.0%	1.1%	0.0%	1.6%
connected_nodes	0.0%	0.8%	0.0%	3.5%
edge_existence	0.0%	0.2%	0.1%	0.8%
node_count	0.0%	2.5%	0.0%	0.1%

Table 6: Share of main-sweep generations that reach the 8,192-token budget, by task and arm, pooled over condition and density. **Bold:** above 50%. In the high-density extension no (task, arm) exceeds 2.5%. All but one of the 78,538 generations that terminate yield an answer.