

Raw accuracy hides a collapse: at $p \geq 0.65$ the model answers 'yes' to 98% of pairs

