

Raw accuracy hides a collapse: at  $p \geq 0.65$  the model answers 'yes' to 98% of pairs

